

Operating Systems (CSE531)

Lecture # 13



Manish Shrivastava

LTRC, IIIT Hyderabad

Process Synchronization

Manish Shrivastava

Race Condition

- **Race condition:** The situation where several processes access – and manipulate shared data concurrently. The final value of the shared data depends upon which process finishes last.
- To prevent race conditions, concurrent processes must be **synchronized**.

Critical Section Problem

- n processes all competing to use some shared data
- Each process has a code segment, called *critical section*, in which the shared data is accessed.
- Problem – ensure that when one process is executing in its critical section, no other process is allowed to execute in its critical section.

Background

- Concurrent access to shared data may result in data inconsistency.
- Maintaining data consistency requires mechanisms to ensure the orderly execution of cooperating processes.
 - Shared-memory solution to bounded-buffer problem allows at most $n - 1$ items in buffer at the same time. A solution, where all N buffers are used is not simple.
 - Suppose that we modify the producer-consumer code by adding a variable *counter*, initialized to 0 and incremented each time a new item is added to the buffer

Bounded-Buffer

- Shared data

```
#define BUFFER_SIZE 10
typedef struct {
    . . .
} item;
item buffer[BUFFER_SIZE];
int in = 0;
int out = 0;
int counter = 0;
```

Bounded-Buffer

- Producer process

item nextProduced;

```
while (1) {  
    while (counter == BUFFER_SIZE)  
        ; /* do nothing */  
    buffer[in] = nextProduced;  
    in = (in + 1) % BUFFER_SIZE;  
    counter++;  
}
```

Bounded-Buffer

- Consumer process

item nextConsumed;

```
while (1) {  
    while (counter == 0)  
        ; /* do nothing */  
    nextConsumed = buffer[out];  
    out = (out + 1) % BUFFER_SIZE;  
    counter--;  
}
```


Bounded Buffer

- Although, both the producer and consumer routines are correct separately they may not function correctly when executed concurrently.
- The statements

counter++;
counter--;

must be performed *atomically*.

- Atomic operation means an operation that completes in its entirety without interruption.

Bounded Buffer

- The statement “**counter++**” may be implemented in machine language as:

register1 = counter

register1 = register1 + 1

counter = register1

- The statement “**count—**” may be implemented as:

register2 = counter

register2 = register2 – 1

counter = register2

Bounded Buffer

- If both the producer and consumer attempt to update the buffer concurrently, the assembly language statements may get interleaved.
- Interleaving depends upon how the producer and consumer processes are scheduled.

Bounded Buffer

- Assume **counter** is initially 5. One interleaving of statements is:

producer: **register1** = **counter** (*register1* = 5)

producer: **register1** = **register1** + 1 (*register1* = 6)

consumer: **register2** = **counter** (*register2* = 5)

consumer: **register2** = **register2** - 1 (*register2* = 4)

producer: **counter** = **register1** (*counter* = 6)

consumer: **counter** = **register2** (*counter* = 4)

- The value of **count** may be either 4 or 6, where the correct result should be 5.

Solution

- A solution to critical section problem must satisfy the following conditions.
 - **Mutual Exclusion.** If process P_i is executing in its critical section, then no other processes can be executing in their critical section.
 - **Progress.** At least one process requesting entry into CS will be able to enter it if there is no other process in it..
 - **Bounded Waiting.** No process waits indefinitely to enter CS once it has requested entry.
 - Assume that each process executes at a nonzero speed
 - No assumption concerning relative speed of the n processes.

Approaches

- Several kernel-level processes may active at a time
 - Example: Data structure “List of open files”
- Kernel developers should ensure that OS is free from race conditions.
- Two approaches are used

Non-preemptive

- Non-preemptive kernel
 - A non-preemptive kernel does not allow a process running in the kernel mode to be preempted.
 - Kernel mode process runs until it exists kernel mode, blocks, or voluntarily yields the control of CPU
 - Free from race conditions

Preemptive

- Preemptive kernel
 - A preemptive kernel allows a process to be pre-empted while it is running in kernel mode.
 - Should be carefully designed
 - Difficult to design especially in SMP
- Why we prefer preemptive kernels ?
 - Suitable for real-time programming
 - More responsive as kernel mode process can not run for a longer time.
- WINDOWS XP, WINDOWS 2000, Prior to LINUX 2.6 are non-preemptive
- Solaris and IRIX are preemptive

Mutual exclusion: Software approaches

- Software approaches can be implemented
- Assume elementary mutual exclusion at the memory access level.
 - Simultaneous access to the same location in main memory are serialized in some order.
- Beyond this, no other support in the hardware, OS, programming language is assumed.

Two process solution : Dekker's algorithm

- ▶ Reported by Dijkstra, 1965.
- ▶ Only 2 processes, P_0 and P_1
- ▶ General structure of process P_i (other process P_j)
 do {
 entry section
 critical section
 exit section
 reminder section
 } **while (1);**
- ▶ Processes may share some common variables to synchronize their actions.

Algorithm 1

- Shared variables:
 - **int** **turn**;
initially **turn = 0**
- Turn variable

P0 while (turn != 0) ; <i>/* Do nothing */</i> <i>critical section</i> turn = 1; remainder section	P1 while (turn != 1); <i>/* Do nothing */</i> <i>critical section</i> turn = 0; remainder section
--	---

- Shared variable *turn* indicates who is allowed to enter next, can enter if *turn = me*
- On exit, point variable to other process
- Deadlock if other process never enters

- +Satisfies mutual exclusion: Only one process can enter in CS
- -It does not satisfy the progress requirement, as it requires strict alternation of processes to enter CS.
- The pace of execution is dictated by slower process.
- If $turn=0$, P1 is ready to enter into CS, P1 can not do so, even though P0 may be in the RS.
- If one process fails in CS or RS, other process is blocked permanently.

Algorithm 2

▶ Problem with Alg1

- It does not retain sufficient information about the state of each process.
- Alg1 remembers only which process is allowed to enter the CS.

▶ To solve this problem, variable turn is replaced by **boolean flag[2]**; flag[0] is for P0; and flag[1] is for P1.

- ▶ Each process may examine the other's flag but may not alter it.
- ▶ When a process wishes to enter CS, it periodically checks other's flag until that flag is false (other process is not in CS)
- ▶ The process sets its own flag true and enters CS.
- ▶ When it leaves CS, it sets its flag to false.

Algorithm 2...

► initially **flag [0] = flag [1] = false.**

► **P0**

```
while ( flag[1] ) ;  
/* Do nothing */  
flag[0] = true;  
critical section  
flag[0] = false;
```

P1

```
while ( flag[0] )  
/* Do nothing */  
flag[1] = true;  
critical section  
flag[1] = false;
```

► Mutual exclusion is satisfied.

► If one process fails outside CS the other process is not blocked.

► Sometimes, the solution is worst than previous solution.

- It does not even **guarantee ME**.
 - P0 executes the **while** statement and finds flag[1] set to false.
 - P1 executes the **while** statement and finds flag[0] set to false.
 - P0 sets flag[0] to true and enters its CS.
 - P1 sets flag[1] to true and enters its CS.

Algorithm 3

- ▶ Interchange the first two statements.
- ▶ Busy Flag Modified

P0	P1
<pre>flag[0] = true; while (flag[1]); <i>/* Do nothing */</i> <i>critical section</i> flag[0] = false;</pre>	<pre>flag[1] = true; while (flag[0]); <i>/* Do nothing */</i> <i>critical section</i> flag[1] = false;</pre>

- ▶ Guarantees ME
- ▶ Both processes set their flags to true before either has executed the **while** statement, then each will think the other has entered CS causing deadlock.

Correct solution (1)

- ▶ Combining the key ideas of previous algorithms
- ▶ Dekker's Algorithm
 - Use *flags* for mutual exclusion, *turn* variable to break deadlock
 - Handles mutual exclusion, deadlock, and starvation

Dekker's Algorithm

- Initial state: $\text{flag}[0]=\text{flag}[1]=\text{false}; \text{turn}=1$

P0

```
flag[0] = true;
while ( flag[1] )
    if (turn==1)
    {
        flag[0]=false;
        while (turn==1)
            /* do nothing */
        flag[0]=true;
    }
/* critical section */
turn=1;
flag[0] = false;
remainder section
```

P1

```
flag[1] = true;
while ( flag[0])
    if (turn==0)
    {
        flag[1]=false;
        while (turn==0)
            /* do nothing */
        flag[1]=true;
    }
/* critical section */
turn=0;
flag[1] = false;
remainder section
```

Correct solution (2)

- Peterson's Algorithm
- Initial state: flag[0]=flag[1]=false;

P0

flag[0] = true;

turn = 1;

while (flag[1] && turn==1)

 /* Do Nothing */;

critical section

flag[0] = false;

remainder section

P1

flag[1] = true;

turn = 0;

while (flag[0] && turn==0)

 /* Do nothing */;

critical section

flag[1] = false;

remainder section

Correct solution

- We need to show that
 - ME is preserved
 - The progress requirement is satisfied
 - The bounded-waiting requirement is met.
- **ME is preserved**
 - If both processes enter the CS both `flag[0]==flag[1]==true`
 - Both could not execute while loop successfully as turn is either 0 or 1.
- **Progress.**
 - While P1 exits CS it sets `flag[1]=false`, allowing P0 to enter CS.
 - P1 and P0 will enter the CS (Progress)
- **Bounded waiting:** P1 will enter the CS after at most one entry by P0 and vice versa.

Multi-process solution: Bakery Algorithm

- ▶ Based on scheduling algorithm commonly used in bakeries.
 - On entering the store the customer receives the number.
 - The customer with the lowest number is served.
 - Customers may receive the same number, then the process with the lowest name is served first.
- ▶ Before entering its critical section, process receives a number. Holder of the smallest number enters the critical section.
- ▶ If processes P_i and P_j receive the same number, if $i < j$, then P_i is served first; else P_j is served first.
- ▶ The numbering scheme always generates numbers in increasing order of enumeration; i.e.,
1,2,3,3,3,3,4,5...

Bakery Algorithm

- var: choosing: array[0...n-1] of boolean.
- Notation \leq lexicographical order (ticket #, process id #)
 - $(a,b) < c,d$ if $a < c$ or if $a = c$ and $b < d$
 - $\max(a_0, \dots, a_{n-1})$ is a number, k , such that $k \geq a_i$ for $i=0, \dots, n-1$
- Shared data

boolean choosing[n];

int number[n];

Data structures are initialized to **false** and **0** respectively

Bakery Algorithm

```
do {  
    choosing[i] = true;  
    number[i] = max(number[0], number[1], ..., number [n – 1])+1;  
    choosing[i] = false;  
    for (j = 0; j < n; j++) {  
        while (choosing[j]) ;  
        while ((number[j] != 0) && (number[j,j] < number[i,i])) ;  
    }  
    critical section  
    number[i] = 0;  
    remainder section  
} while (1);
```

- **Consider P_i in its CS and P_k is trying to enter CS**
- When P_k enters second while statement for $j=i$, it finds that
 - $\text{number}[i] \neq 0$
 - $(\text{number}[i], i) < (\text{number}[k], k)$
 - So it leaves until P_i leaves CS
- FCFS is followed.

Mutual exclusion: hardware solution

- ▶ In the uni-processor system, it is sufficient to prevent a process from being interrupted.

```
while (true){  
    /* disable interrupts */  
    /* Critical section */  
    /* enable interrupts */  
    /* remainder */  
}
```

- ▶ **Since CS can not be interrupted ME is guaranteed.**
- ▶ **The efficiency decreases**
- ▶ **It can not work in multi-processor environments**
 - More than one process is executing at a time.

Special machine instructions

- In multi-processor configuration, several processes share access to a common main memory.
- At the hardware level, access to a memory location excludes any other access to that same memory location.
- Processor designers have proposed several machine instructions to carry out two actions atomically (single cycle).
 - Reading and writing
 - swapping

Test and set instruction

- Test and modify the content of a word atomically

```
boolean testset (int i)
{
  if (i==0)
  {
    i=1;
    return true;
  }
  else
  {
    return false
  }
}
```

- This instruction sets the value of 'i', if the value=0 and returns true. Otherwise the value is not changed and false is returned.

Mutual Exclusion with Test-and-Set

- ▶ Shared data:

```
    boolean lock = false;
```

- ▶ void *P*(int *i*)

```
    do {
```

```
        while (TestAndSet(lock)==false)
```

```
            /* do nothing*/;
```

```
            critical section
```

```
            lock = false;
```

```
            remainder section
```

```
    }
```

Test-and-Set: Correctness

■ Mutual exclusion

- ✦ A shared variable lock is set to false
- ✦ The only process P_i that enters CS that finds lock as false and sets it to true.
- ✦ All other processes trying to enter CS so into a busy waiting mode and finds lock as false.
- ✦ When process leaves C it resets lock to false.
- ✦ When P_i exits lock is set to false so the next process P_j to execute instruction find test-and-set=false and will enter the CS.

■ Progress

- ✦ Trivially true

■ Unbounded waiting

- ✦ Possible since depending on the timing of evaluating the test-and-set primitive.
- ✦ Does not guarantee fairness.

Swap instruction

- Atomically swap two variables.

```
void swap(boolean &a, boolean &b) {  
    boolean temp = a;  
    a = b;  
    b = temp;  
}
```

Mutual Exclusion with Swap

- Shared data (initialized to **false**):
 boolean lock;
 boolean waiting[n];
- Process P_i
 do {
 key = true;
 while (key == true)
 Swap(lock,key);
 critical section
 lock = false;
 remainder section
 }

SWAP: Correctness

- Similar to Test-and-set
- Mutual exclusion
- Progress
 - ✦ Trivially true
- Unbounded waiting
 - ✦ Possible since depending on the timing of evaluating the test-and-set primitive.
 - ✦ Does not guarantee fairness.

Can we get bounded waiting ?

- Introduce a boolean array called waiting of size n and boolean variable key

- Entry

- ✦ waiting[i]:=true;
- ✦ key:=true;
- ✦ while (waiting[i] and key) do
 - ✓ Swap(&key,&lock)
- ✦ waiting[i]:=false;
- ✦ execute CRITICAL SECTION

- Exit

- ✦ Find the next process j that has waiting[j]=1 stepping through waiting.
- ✦ Set waiting[j]:=false;
- ✦ Process P_j immediately enter the CS.
- ✦ If no process exists, set lock=false;

Can we get bounded waiting ?....

- Every (interested) P_i executes test&set at least once.
- P_i enters the critical section provided:
 - ✦ Key is false in which case there is no process in CS.
- Or
 - ✦ If it was waiting, because `waiting[i]` was reset to false by the unique process that was blocking it in the critical section.
 - ✦ Either of the above events occur exactly once and hence mutual exclusion.

Properties of machine instruction approach

- +ve
 - Any number of processes
 - Simple and easy
 - Can support multiple CSs.
- -ve
 - Busy waiting is employed
 - The process is waiting and consuming processor time.
 - Starvation is possible.
 - The selection of waiting process is arbitrary.
 - Deadlock is possible due to priority
 - P1 enters CS and interrupted by higher priority process P2 which is trying to enter CS.
 - P2 can not get CS unless P1 is out and P1 can not be dispatched due to low priority.

Semaphores: Dijkstra; 1965

- ▶ Two and more processes can cooperate by means of simple signals, such that a process is forced to stop at a specified place until it has received a specific signal.
- ▶ For signaling, special variables called semaphores are used
- ▶ A semaphore is a synchronization tool.
- ▶ A semaphore is an integer variable that is accessed only through two standard atomic operations: **wait and signal**.
- ▶ To transmit a signal to semaphore S, a process executes the primitive *signal(S)* primitive.
- ▶ To receive a signal via semaphore S, the process executes *wait(S)* primitive.

Semaphores: Dijkstra 1965

Classical or first definition

- A semaphore is initialized to a non-negative value
- The **wait** operation decrements the semaphore value. If the integer value is negative the process waits.
- The **signal** operation increments the semaphore value. If the value is not positive, then process which is blocked by a wait operation is gets the access to CS.
- The wait and signal are assumed to be atomic.
- Semaphore S – integer variable
- can only be accessed via two indivisible (atomic) operations

wait (S):

while $S \leq 0$ do *no-op*;
 $S--$;

signal (S):

$S++$;

Critical Section of n Processes

- ▶ Shared data:

semaphore mutex; // initially $mutex = 1$

- ▶ Process P_i :

```
do {  
    wait(mutex);  
    critical section  
    signal(mutex);  
    remainder section  
} while (1);
```

- ▶ **Modifications to the integer value of the semaphore in the wait and signal operations must be executed indivisibly.**

Semaphore Implementation

- The classical definition requires busy waiting.
- While a process is in CS, the other process must loop continuously in the entry code.
- Busy waiting wastes CPU cycles.
- This type of semaphore is called spinlock: process spins while waiting for a lock.
 - Advantage of spinlock: no context switch
 - When locks are expected to be held for short times, spinlocks are useful.
- To overcome the need for busy waiting, we can modify the definition of the wait and signal semaphore operations.
- If a process executes wait operation and finds the semaphore operation is not positive, it must wait.
 - Rather than busy waiting it must block itself.
 - The **block** operation puts the process into waiting queue of semaphore and process is switched to waiting state.
- A process that is blocked waiting on a semaphore S, should be restarted when some other process executes signal operation.
- The process is restarted with **wakeup** operation.

Semaphore Implementation

- Define a semaphore as a record

```
typedef struct {  
    int value;  
    struct process *L;  
} semaphore;
```

- Assume two simple operations:
 - **block** suspends the process that invokes it.
 - **wakeup(*P*)** resumes the execution of a blocked process **P**.

Implementation

- Semaphore operations now defined as

wait(S):

```
S.value--;  
if (S.value < 0) {  
    add this process to S.L;  
    block;  
}
```

signal(S):

```
S.value++;  
if (S.value <= 0) {  
    remove a process P from S.L;  
    wakeup(P);  
}
```

- Wait and signal operations are system calls.

Semaphore as a General Synchronization Tool

- Execute B in P_j only after A executed in P_i
- Use semaphore $flag$ initialized to 0
- Code:

P_i	P_j
\vdots	\vdots
A	$wait(flag)$
$signal(flag)$	B

Deadlock and Starvation

- **Deadlock** – two or more processes are waiting indefinitely for an event that can be caused by only one of the waiting processes.
- Let S and Q be two semaphores initialized to 1

P_0	P_1
<i>wait(S);</i>	<i>wait(Q);</i>
<i>wait(Q);</i>	<i>wait(S);</i>
\vdots	\vdots
<i>signal(S);</i>	<i>signal(Q);</i>
<i>signal(Q)</i>	<i>signal(S);</i>

- **Starvation** – indefinite blocking. A process may never be removed from the semaphore queue in which it is suspended.

Two Types of Semaphores

- *Counting* semaphore – integer value can range over an unrestricted domain.
- *Binary* semaphore – integer value can range only between 0 and 1; can be simpler to implement.

Binary Semaphores

- ▶ A binary semaphore is a semaphore with an integer value that can range only between 0 and 1

- ▶ It is simple to implement.

- ▶ Type binary semaphore = **record**

```
    value:(0,1)  
    queue: list of processes
```

```
end;
```

- ▶ var s: binary semaphore

- ▶ **waitB(s):**

```
    If s.value=1 then
```

```
        s.value=0
```

```
    else
```

```
        begin
```

```
            place this process in s.queue;
```

```
            block this process;
```

```
        end;
```

- **signalB(s):**

```
    If s.queue is empty then
```

```
        s.value=1
```

```
    else
```

```
        begin
```

```
            remove a process from s.queue;
```

```
            place this process in the ready list.
```

```
        end;
```

Implementing S as a Binary Semaphore

- Can implement a counting semaphore S as a binary semaphore.
- Data structures:
binary-semaphore S1, S2;
int C;
- Initialization:
S1 = 1
S2 = 0
C = initial value of semaphore S

- *wait* operation
wait(S1);
C--;
if (C < 0) {
 signal(S1);
 wait(S2);
 }
signal(S1);
- *signal* operation
wait(S1);
C ++;
if (C <= 0)
 signal(S2);
else
 signal(S1);

Counting semaphores

wait(S):
 S.value--;
 if (S.value < 0) {
 add this process to S.L;
 block;
 }

signal(S):
 S.value++;
 if (S.value <= 0) {
 remove a process P from
 S.L;
 wakeup(P);
 }

Implementing wait() and signal() in Multi-processor Systems

- Disabling interrupts will not work.
- Spinlock is the solution
- With this we have moved busy waiting from entry section to critical sections of application programs.