

[illegible]

Manish Shrivastava

LTRC, IIIT Hyderabad

Classical Problems of Synchronization

- Bounded-Buffer Problem
- Readers and Writers Problem
- Dining-Philosophers Problem

Bounded-Buffer Problem

- ▶ Used to to illustrate the power of synchronization techniques
- ▶ We assume that the buffer consists of n buffers, each capable of holding an item.
- ▶ The mutex semaphore provides mutual exclusion access to buffer which is initialized to the value 1.
- ▶ The **empty** and **full** semaphores count the number of empty and full buffers which are initialized to n and zero respectively.
- ▶ Shared data
semaphore full, empty, mutex;
- ▶ Initially:
full = 0, empty = n , mutex = 1

Bounded-Buffer Problem Producer Process

```
do {  
    ...  
    produce an item in nextp  
    ...  
    wait(empty);  
    wait(mutex);  
    ...  
    add nextp to buffer  
    ...  
    signal(mutex);  
    signal(full);  
} while (1);
```

Bounded-Buffer Problem Consumer Process

```
do {  
    wait(full)  
    wait(mutex);  
    ...  
    remove an item from buffer to nextc  
    ...  
    signal(mutex);  
    signal(empty);  
    ...  
    consume the item in nextc  
    ...  
} while (1);
```

- Producer is producing full buffers for the consumer and consumer is producing empty buffers for the consumer.

Readers-Writers Problem

- ▶ Problem: A data object (file or record) is shared among several concurrent processes.
 - Some want to read and others want to update it.
- ▶ **Readers:** processes interested in reading.
- ▶ **Writers:** processes interested in writing.
- ▶ Two readers can access shared data object simultaneously.
- ▶ But a writer and reader can access shared data object simultaneously
 - problems may occur!
- ▶ To protect from these problems, writers should have an exclusive access to the shared object.
- ▶ This synchronization problem is referred to as readers-writers problem.
- ▶ It is a different kind of synchronization problem.
- ▶ The readers-writers problem has several variations.
 - Simple one: No reader will be kept waiting unless writer has obtained permission to write.

Readers-Writers Problem

- ▶ Semaphores used: **mutex** and **wrt**
- ▶ The semaphore **wrt** is common to reader and writer.
- ▶ Semaphore **mutex** is used to update **readcount**.
- ▶ **readcount** keeps track of how many are reading the object.

- ▶ Shared data

semaphore mutex, wrt;

Initially

mutex = 1, wrt = 1, readcount = 0

Readers-Writers Problem Writer Process

wait(wrt);

...

writing is performed

...

signal(wrt);

Readers-Writers Problem Reader Process

```
wait(mutex);
readcount++;
if (readcount == 1)
    wait(wrt);
signal(mutex);
...
reading is performed
...
wait(mutex);
readcount--;
if (readcount == 0)
    signal(wrt);
signal(mutex);
```


- Writers can be starved if there is a continuous sequence of readers.

Readers-Writers Problem

- Can the producer/consumer problem is considered as case of the readers/writers problem with a writer is a producer and reader is a consumer ?
- The answer is no
- The producer is not just a writer
 - It must read queue pointers to determine where to write the next item and it must determine if the buffer is full.
- Similarly the consumer is not a reader
 - It must adjust queue pointers to show that it has removed a unit from the buffer.

Dining-Philosophers Problem

- ▶ Five philosophers spend their lives thinking and eating.
- ▶ They share a common circular table surrounded by five chairs.
- ▶ Five single chopsticks are available.
- ▶ Whenever a philosopher wants to eat, he tries to pick up two chopsticks that are closest to him/her.
- ▶ A philosopher can not pick the chopstick in the hand of neighbor.
- ▶ After finishing, the philosopher puts back the chopsticks and starts thinking.
- ▶ It is simple representation of the need to allocate several resources among several processes in a **deadlock and starvation free manner**.



Dining-Philosophers Problem

- Shared data

semaphore chopstick[5];

Initially all values are 1

- Philosopher i :

```
do {  
    wait(chopstick[i])  
    wait(chopstick[(i+1) % 5])  
    ...  
    eat  
    ...  
    signal(chopstick[i]);  
    signal(chopstick[(i+1) % 5]);  
    ...  
    think  
    ...  
} while (1);
```

- The solution creates a deadlock

Barbershop Problem

- 3 barbers, each with a barber chair
 - Haircuts may take varying amounts of time
- Sofa can hold 4 customers, max of 20 in shop
- Customers wait outside if necessary
- When a chair is empty:
 - Customer sitting longest on sofa is served
 - Customer standing the longest sits down
- After haircut, go to cashier for payment
 - Only one cash register
 - Algorithm has a separate cashier, but often barbers also take payment
 - This is also a critical section

Barbershop Problem

- The main body of the program activates 50 customers, 3 barbers, and the cashier process. Synchronization operators.
 - Shop and sofa capacity: the capacity of shop and the capacity of the sofa are governed by the semaphores **max_capacity** and **sofa**.
 - When customer enters max_capacity decremented by one.
 - When a customer leaves it is incremented.
 - Wait and signal operations are surround the actions of sitting and getting_up from sofa.
 - **Barber chair capacity:**
 - There are three barber chairs; the semaphore barber_chair assures that no more than three customers attempt to obtain service at a time.
 - A customer will not get up from the sofa until at least one chair is free.
 - **Ensuring customers are in the barber chair:** The semaphore cust_ready provides a wakeup signal for a sleeping barber indicating that the customer has just taken the chair.
 - **Holding customers in barber chair:** once seated the customer remain in the chair until the barber gives the signal that haircut is complete, using the semaphore finished.
 - **Limiting one customer to a barber chair:** the semaphore barber_chair is intended to limit the number of customers in barber chairs to three. The semaphore leave_b_chair is used to synchronize sitting.
 - **Paying and receiving:** payment and receipt semaphores are used to synchronize the operations.
 - **Coordinating barber and cashier functions:** To save money the barber shop does not employ a separate cashier. Each barber is required to perform that task when not cutting hair. The semaphore coord ensures the barbers perform only one task at a time.

Barbershop Problem

| Semaphore | Wait operation | Signal operation |
|----------------------|--|---|
| max_capacity | Customer waits for a room to enter shop. | Exiting customer signals customer waiting to enter |
| sofa | Customer waits for seat on sofa | Customer leaving sofa signals customer waiting for sofa |
| barber_chair | Customer waits for empty barber chair | Barber signals when that barber's chair is empty |
| Cust_read | Barber waits until customer is in the chair | Customer signals barber that customer is in the chair |
| finished | Customer waits until his haircut is complete. | Barber signals when done cutting hair of his customer. |
| leave_b_chair | Barber waits until customer gets up from the chair | Customer signals barber when customer gets up from chair. |
| payment | Cashier waits for a customer to pay | Customer signals cashier that he has paid. |
| receipt | Customer waits for a receipt for a payment | Cashier signals that payment has been accepted. |
| coord | Wait for a barber resource to be free to be free perform either the hair cutting or cashiering function. | Signal that a barber resource is free. |

Barbershop

```
program barbershop1;
var
  max_capacity: semaphore (:=20);
  sofa: semaphore (:=4);
  barber_chair, coord: semaphore (:=3);
  cust_ready, leave_b_chair, payment, receipt: semaphore (:=0)
```

```
procedure customer;
var custnr: integer;
begin
  wait (max_capacity );
  enter shop;

  wait( sofa );
  sit on sofa;
  wait( barber_chair );
  get up from sofa;
  signal( sofa );
  sit in barber chair;
  wait( mutex2 );
  signal( cust_ready );
  wait( finished[custnr] );
  leave barber chair;
  signal( leave_b_chair );
  pay;
  signal( payment );
  wait( receipt );
  exit shop;
```

```
procedure barber;
var b_cust: integer
begin
  repeat
    wait( cust_ready );

    wait( coord );
    cut hair;
    signal( coord );
    signal( finsihed[b_cust] );
    wait( leave_b_chair );
    signal( barber_chair );
  forever
end;
```

```
procedure cashier;
begin
  repeat
    wait( payment );
    wait( coord );
    accept payment;
    signal( coord );
    signal( receipt );
  forever
end;
```

```
Void main()
{
  count=0;
  Parbegin {customer... 50 times,...customer,
  Barber, barber,barber, cashier)
}
```


Barbershop Problem

- The preceding solution is unfair.
- The customers are served in the order they enter the shop.
- If one barber is very fast and one of the customer is quite bald.
- The problem can be solved with more semaphores.
 - We assign unique customer number to each customer.
 - The semaphore mutex1 protects access to global variable count.
- The semaphore finished is refined to be an array of 50 semaphores.
 - Once a customer seated in a barber chair, he executes `wait(finished[custnt])` to wait in his own unique semaphore.
- Please see the solution in William Stallings book (pp 229-234)

Fair Barbershop

```
program
var
    barbershop2;
    max_capacity: semaphore (:=20);
    sofa: semaphore (:=4);
    barber_chair, coord: semaphore (:=3);
    mutex1, mutex2: semaphore (:=1);
    cust_ready, leave_b_chair, payment, receipt: semaphore (:=0);
    finished: array [1..50] of semaphore (:=0);
    count: integer;
```

```
procedure customer;
var custnr: integer;
begin
    wait (max_capacity );
    enter shop;
    wait( mutex1 );
    count := count + 1;
    custnr := count;
    signal( mutex1 );
    wait( sofa );
    sit on sofa;
    wait( barber_chair );
    get up from sofa;
    signal( sofa );
    sit in barber chair;
    wait( mutex2 );
    enqueue1( custnr );
    signal( cust_ready );
    signal( mutex2 );
    wait( finished[custnr] );
    leave barber chair;
    signal( leave_b_chair );
    pay;
```

```
procedure barber;
var b_cust: integer;
begin
    repeat
        wait( cust_ready );
        wait( mutex2 );
        dequeue1( b_cust );
        signal( mutex2 );
        wait( coord );
        cut hair;
        signal( coord );
        signal( finished[b_cust] );
        wait( leave_b_chair );
        signal( barber_chair );
    forever
end;
```

```
procedure cashier;
begin
    repeat
        wait( payment );
        wait( coord );
        accept payment;
        signal( coord );
        signal( receipt );
    forever
end;
```

```
Void main()
{
    count=0;
    Parbegin {customer... 50 times,...customer,
    Barber, barber,barber, cashier)
}
```

Implementation of Semaphores

- ▶ wait and signal operations are atomic.
- ▶ Good Solution: implement through hardware or firmware.
- ▶ Other solutions
 - Ensure that only process manipulates “wait” and “signal” operations.
 - One can use Dekker’s algorithm or Peterson’s algorithm
 - Substantial processing overhead
 - Use one of the hardware supported schemes
 - Test and set
 - disabling interrupts (single processor)
- ▶ The wait and signal code is very short the amount of busy waiting involved is short.

Two possible implementations of Semaphores

```
Wait(s)
{
    while(!testset(s.flag)
        /* do nothing */
        s.count--;
        if (s.count < 0)
        {
            place this process in s.queue;
            block this process (set s.flag to 0)
        }
    else
        s.flag=0;
}
```

```
Signal(s)
{
    while(!testset(s.flag)
        /* do nothing */
        s.count++;
        if (s.count <= 0)
        {
            remove a process P from s.queue;
            Place a process P in the ready list
        }
        s.flag=0;
}
```

With TestSet Instruction

```
Wait(s)
{
    Inhibit interrupts
    s.count--;
    if (s.count < 0)
    {
        place this process in s.queue;
        block this process allow interrupts
    }
    else
        allow interrupts;
}
```

```
Signal(s)
{
    Inhibit interrupts;
    s.count++;
    if (s.count <= 0)
    {
        remove a process P from s.queue;
        Place a process P in the ready list
    }
    allow interrupts;
}
```

With Interrupts

Problem with semaphores

- Incorrect use may result in timing errors
- These errors are difficult to detect as these occur if only particular sequence occurs.
- Missing or reverse order.
- It is difficult to produce correct program using semaphores.
- The wait and signal operations are scattered throughout the program and it is difficult to see overall effect of these operations on the semaphores.

Problem with semaphores (cont.)

► Problems

- Suppose a process interchanges the order in which wait and signal operations on the semaphore are executed

signal (mutex)
CS
wait (mutex)

- Several processes may be executing in their CS simultaneously.
- Suppose that a process replaces signal(mutex) with wait(mutex)

| | |
|----------------|-------------|
| signal (mutex) | wait(mutex) |
| CS | CS |
| signal (mutex) | wait(mutex) |

- Deadlock will occur.
- Suppose a process omits wait(mutex) or signal(mutex) or both.
 - ME is violated or deadlock occurs.

► A critical region and monitor concept is introduced to address this problem

THANK YOU